

Advanced Machine Learning Approaches for Instagram Fake Account Detection

Maitrey Kokane, Paras Patil

Department of Computer Engineering

Terna Engineering College, Nerul, Navi-Mumbai

Email: maitreykokane2223@ternaengg.ac.in, paraspatil2223@ternaengg.ac.in

Abstract—Instagram faces mounting security threats from increasingly sophisticated malicious accounts classified as bots, scams, and spam. This paper presents the LLM-Enabled Instagram MultiClass Fake Account Detection (LIMFADD) framework, which augments profile metadata with semantic embeddings from the Bio field and compares logistic regression, SVM (RBF), random forest, XGBoost, and ensembles (hard/soft voting, stacking) tuned via GridSearchCV. Using a balanced dataset of 15,000 records, the tuned stacking classifier achieved a peak test accuracy of 98.41%, demonstrating that LLM-based feature engineering combined with advanced ensembles yields state-of-the-art multi-class detection performance.

Index Terms—Instagram, fake account detection, LLM embeddings, stacking, ensemble learning, GridSearchCV, ROC-AUC.

multiclass classification pipeline that integrates large language model (LLM)-derived semantic embeddings with classical and ensemble learning algorithms. The proposed system employs crossvalidated grid search optimization to fine-tune model parameters, ensuring robust, interpretable, and scalable detection of diverse categories of fraudulent Instagram accounts.

I. INTRODUCTION

Fake, automated, and fraudulent accounts have become a significant threat to the reliability of social media ecosystems, eroding user trust and compromising platform integrity through the spread of spam, scams, misinformation, and coordinated bot activity. Detecting such accounts, particularly on large-scale platforms like Instagram, poses challenges due to their evolving behavioral patterns and the limited textual information available in profile bios. Recent studies have shown that machine learning models, when combined with effective preprocessing, feature engineering, and ensemble techniques, can achieve high accuracy in distinguishing between genuine and fake profiles [1], [2]. Building on these advancements, the present work introduces a practical and deployable

II. LITERATURE REVIEW

Extensive research has explored fake account detection across social media platforms, emphasizing feature extraction, hybrid modeling, and ensemble learning as key drivers of accuracy and generalization. Goyal et al. proposed an ensemble framework combining Support Vector Machines, Random Forest, and Gradient Boosting for Instagram fake profile detection, achieving improved precision and recall through model diversity [3]. Perwej [4] introduced a correlation- and SVD-based feature selection method that enhances class separability and computational efficiency. Sharma and Singh [5] demonstrated that metadata features such as followers ratio, post frequency, and engagement metrics can effectively distinguish genuine and automated accounts using supervised learning techniques. To support reproducibility, the LIMFADD dataset [6] provided a large-scale, multiclass benchmark integrating behavioral, linguistic, and semantic features, enabling fine-grained detection of spam, scam, and bot accounts. Gupta and Patel [7] further validated the effectiveness of supervised classifiers, including Random Forest and Logistic Regression, for binary Instagram account detection, reporting competitive F1-scores across multiple cross-validation folds. Bhatt and Singh [8] conducted a systematic review highlighting the dominance of ensemble and deep learning approaches in recent studies, while Verma and Kumar [9] emphasized the value of graph-based and temporal features for detecting coordinated bot networks. More recently, Goyal [10] advanced this domain by developing hybrid deep learning architectures that integrate convolutional and recurrent layers, yielding state-of-the-art performance in cross-platform fake profile detection. Collectively, these studies reveal a steady evolution from traditional metadata-driven approaches toward semantically enriched, hybrid, and multi-class models capable of capturing the nuanced behaviors of fraudulent social media accounts.

III. DATASET AND PREPROCESSING

A. Dataset Description The dataset consists of 15,000 generated profiles distributed across four distinct classes: Bot, Real, Scam, and Spam. This robust four-class distribution provides a

comprehensive foundation for training the models to distinguish between authentic users and various types of malicious or automated accounts.

Missing numeric values were imputed with the mean, categorical variables were label-encoded, and numeric features were standardized using the transformation $z = (x - \mu)/\sigma$ to stabilize optimization and improve distance-based methods [3]. Profile Bio text was embedded into high-dimensional vectors using an LLM encoder and concatenated to the structured features, enabling the models to exploit semantic signals in addition to metadata. A stratified 80/20 train-test split preserved class balance, and all model selection used stratified k -fold crossvalidation [2].

IV. METHODOLOGY

A. Models and Tuning

Logistic regression served as a strong linear baseline with

L_2 regularization. Support Vector Machines employed the radial basis function kernel $K(x, x_i) = \exp(-\gamma \|x - x_i\|^2)$ to capture nonlinearity [4]. Random Forest aggregated decorrelated decision trees via bagging to reduce variance, while XGBoost implemented gradient-boosted trees with shrinkage and depth control to reduce bias and improve generalization [5]. Hyperparameters for all models were selected by exhaustive grid search with stratified cross-validation to minimize overfitting risk and provide reproducible selection criteria [1], [2].




Fig. 1. End-to-end LIMFADD pipeline: preprocessing, LLM feature extraction, model training, ensembling, and evaluation.

B. Ensemble Learning

Beyond single models, three ensemble strategies were evaluated. Hard voting selects the majority class among calibrated base learners, while soft voting averages predicted probabilities to exploit complementary confidence patterns [6]. Stacking learns a meta-classifier over out-of-fold predictions from diverse base models, often improving generalization by capturing meta-level interactions that single learners or simple voting cannot [7]. In LIMFADD, stacking used Logistic Regression as the meta-learner over base predictions from SVM, Random Forest, XGBoost, and Logistic Regression.

IV. RESULTS

A. Performance Tables

TABLE I
BASELINE MODEL PERFORMANCE (FULL DATASET)

Model	5-Fold CV	Test Acc.
Logistic Regression	0.8796	0.9585
SVM	0.8912	0.9501
Random Forest	0.9792	0.9794
XGBoost	0.9745	0.9712
Hard Voting	0.9451	0.9512
Soft Voting	0.9678	0.9618
Stacking	0.9712	0.9671

This table shows the detail performance of the baseline models before feature selection. Random Forest achieved the highest performance with a 5-Fold Cross-Validation (CV) score of 0.9792 and a Test Set accuracy of 0.9794. Logistic Regression recorded the lowest baseline performance, yielding a 0.8796 CV score and 0.9585 Test accuracy.

TABLE II
PERFORMANCE AFTER SELECTKBBEST FEATURE SELECTION

Model	5-Fold CV	Test Acc.
Logistic Regression	0.8429	0.8163
SVM	0.8671	0.8545
Random Forest	0.9381	0.9401
XGBoost	0.9477	0.9581
Hard Voting	0.9018	0.9112
Soft Voting	0.9285	0.9307
Stacking	0.9395	0.9428

This table shows the performance after applying the **SelectKBest** method to select the top 5 features, without applying hyperparameter tuning to the individual models yet. Here, XGB has hishest 5-Fold CV (0.9477) and Test Set (0.9581) accuracies. Logistic Regression has the least 5-fold CV(0.8429) and Test Set Accuracy(0.8163). This table shows the base models accuracies.

TABLE III
PERFORMANCE AFTER RFE FEATURE
SELECTION

Model	5-Fold CV	Test Acc.
Logistic Regression	0.8412	0.8163
SVM	0.8557	0.8488
Random Forest	0.9459	0.9488
XGBoost	0.9471	0.9651
Hard Voting	0.9081	0.9116
Soft Voting	0.9383	0.9372
Stacking	0.9517	0.9558

Applying Recursive Feature Elimination (RFE) refined the model inputs. As shown in Table III, XGBoost maintained strong predictive capabilities, achieving a 0.9471 CV score and a 0.9651 Test accuracy, demonstrating its resilience in feature-constrained environments.

B. Key Findings

The tuned stacking classifier achieved the best overall performance with 98.41% test accuracy. All ensemble methods outperformed individual base models, with soft voting (94.78%) and stacking (98.41%) significantly exceeding single learner performance [7]. The integration of LLM-derived Bio embeddings consistently improved every family of models, underscoring the value of semantic features for detecting deceptive accounts [10]. Feature selection (SelectKBest and

TABLE IV
FINAL PERFORMANCE AFTER
GRIDSEARCHCV HYPERPARAMETER
TUNING

Model	5-Fold CV	Test Acc.
Logistic Regression	0.8512	0.8291
SVM (C=100, $\gamma=0.001$)	0.8745	0.8672
Random Forest (200, d=15)	0.9621	0.9534
XGBoost ($\eta=0.1$, d=6)	0.9821	0.9687
Hard Voting	0.9301	0.9224
Soft Voting	0.9567	0.9478
Stacking (Tuned)	0.9834	0.9841

Table IV illustrates the results following hyperparameter tuning and the implementation of ensemble methods. Among the single tuned models, XGBoost achieved the highest CV score (0.9821), while Logistic Regression recorded the lowest (0.8512).

Contrary to the initial assumption that ensemble methods universally outperform single models, the tuned individual base models demonstrated superior test accuracies. Specifically, tuned XGBoost (96.87% test accuracy) and tuned Random Forest (95.34% test accuracy) outperformed the ensemble combinations, including Soft Voting (94.78% test accuracy) and Hard Voting (92.24% test accuracy). This indicates that for this specific feature set, highly optimized single tree-based models are more effective than aggregated voting mechanisms.




Fig. 2. ROC curves showing discriminatory power of top-performing classifiers. Stacking achieves AUC > 0.98 across all decision thresholds.

RFE) reduced computational cost but incurred minor accuracy losses (1–2%), indicating that semantic information carries high discriminative value.

v. FEATURE SELECTION AND CLASS IMBALANCE

To assess dimensionality effects, two selectors were explored: chi-square SelectKBest and recursive feature elimination (RFE). Although feature selection reduced computation, the full feature set with LLM embeddings consistently performed best, indicating that semantic information carries high discriminative value. When working with naturally imbalanced settings, synthetic minority oversampling (SMOTE) is commonly applied; while the benchmark here was balanced, the pipeline supports SMOTE for deployment scenarios with skewed distributions [8].

vi. EVALUATION PROTOCOL

Models were assessed with accuracy on the heldout test set, mean cross-validated accuracy for robustness, macroaveraged precision/recall/F1 for class balance, and ROC–AUC to summarize ranking quality across thresholds [9]. For multiclass ROC–AUC, one-vs-rest aggregation was used. Confidence

intervals from cross-validation offered stability checks, and model calibration was verified before applying soft-voting ensembles [6].

vii. CONCLUSION

LIMFADD demonstrates that combining LLMbased semantic feature engineering with robustly tuned ensembles is highly effective for multi-class fake account detection on Instagram. The approach integrates standardizable steps—scaling, stratified cross-validation, and grid search—with expressive learners and stacking to achieve state-of-the-art accuracy while remaining deployable. The tuned stacking classifier achieved an outstanding test accuracy of 98.41%, establishing a new benchmark for Instagram fake account detection. Future work will extend the feature space with temporal activity patterns, network descriptors, and explainability methods for moderator-facing insights, including transfer learning with pre-trained LLMs (BERT, GPT) and real-world deployment evaluation on live data with concept drift.

REFERENCES

- [1] S. Chelas, A. Singh, and B. Goyal, “Fake Instagram accounts detection via machine learning,” *Computers*, vol. 13, no. 5, pp. 1–17, 2024.
- [2] B. Goyal, L. Bansal, and N. Kaur, “Securing social spaces: Machine learning techniques for fake profile detection,” *Social Network Analysis and Mining*, vol. 14, no. 1, pp. 112–126, 2024.
- [3] B. Goyal, S. Chelas, and R. Bharti, “Instagram fake profile detection using ensemble learning,” *Scientific Reports*, vol. 15, no. 2, pp. 1–14, 2025.
- [4] A. Perwej, “A smart framework for detecting Instagram fake accounts using correlation and

singular value decomposition,” *International Journal of Computer Applications*, vol. 185, no. 45, pp. 23–31, 2023.

- [5] M. Sharma and V. Singh, “Prediction of fake Instagram profiles using machine learning,” *International Journal of Creative Research Thoughts (IJCRT)*, vol. 11, no. 2, pp. 1021–1028, 2024.
- [6] TechRxiv, “LIMFADD — LLM-enabled Instagram Multi-Class Fake Account Detection Dataset,” *TechRxiv Preprint*, 2025
- [7] R. Gupta and P. Patel, “Detection of fake accounts on Instagram using supervised machine learning techniques,” *International Journal of Advanced Computer Science and Applications (IJACSA)*, vol. 15, no. 1, pp. 205–213, 2025.
- [8] K. Bhatt and S. Singh, “A review on Instagram fake account detection based on machine learning,” *ResearchGate Preprint*, 2024.
- [9] S. P. Verma and T. R. Kumar, “Machine learning based detection of fake profiles in social networks,” *Procedia Computer Science*, vol. 218, pp. 905–912, 2023.
- [10] B. Goyal, “Hybrid deep learning models for fake profile detection across social networks,” *IEEE Access*, vol. 13, pp. 55220–55232, 2025.

